



动量与学习率衰减

主讲人：龙良曲

Tricks

- momentum
- learning rate decay



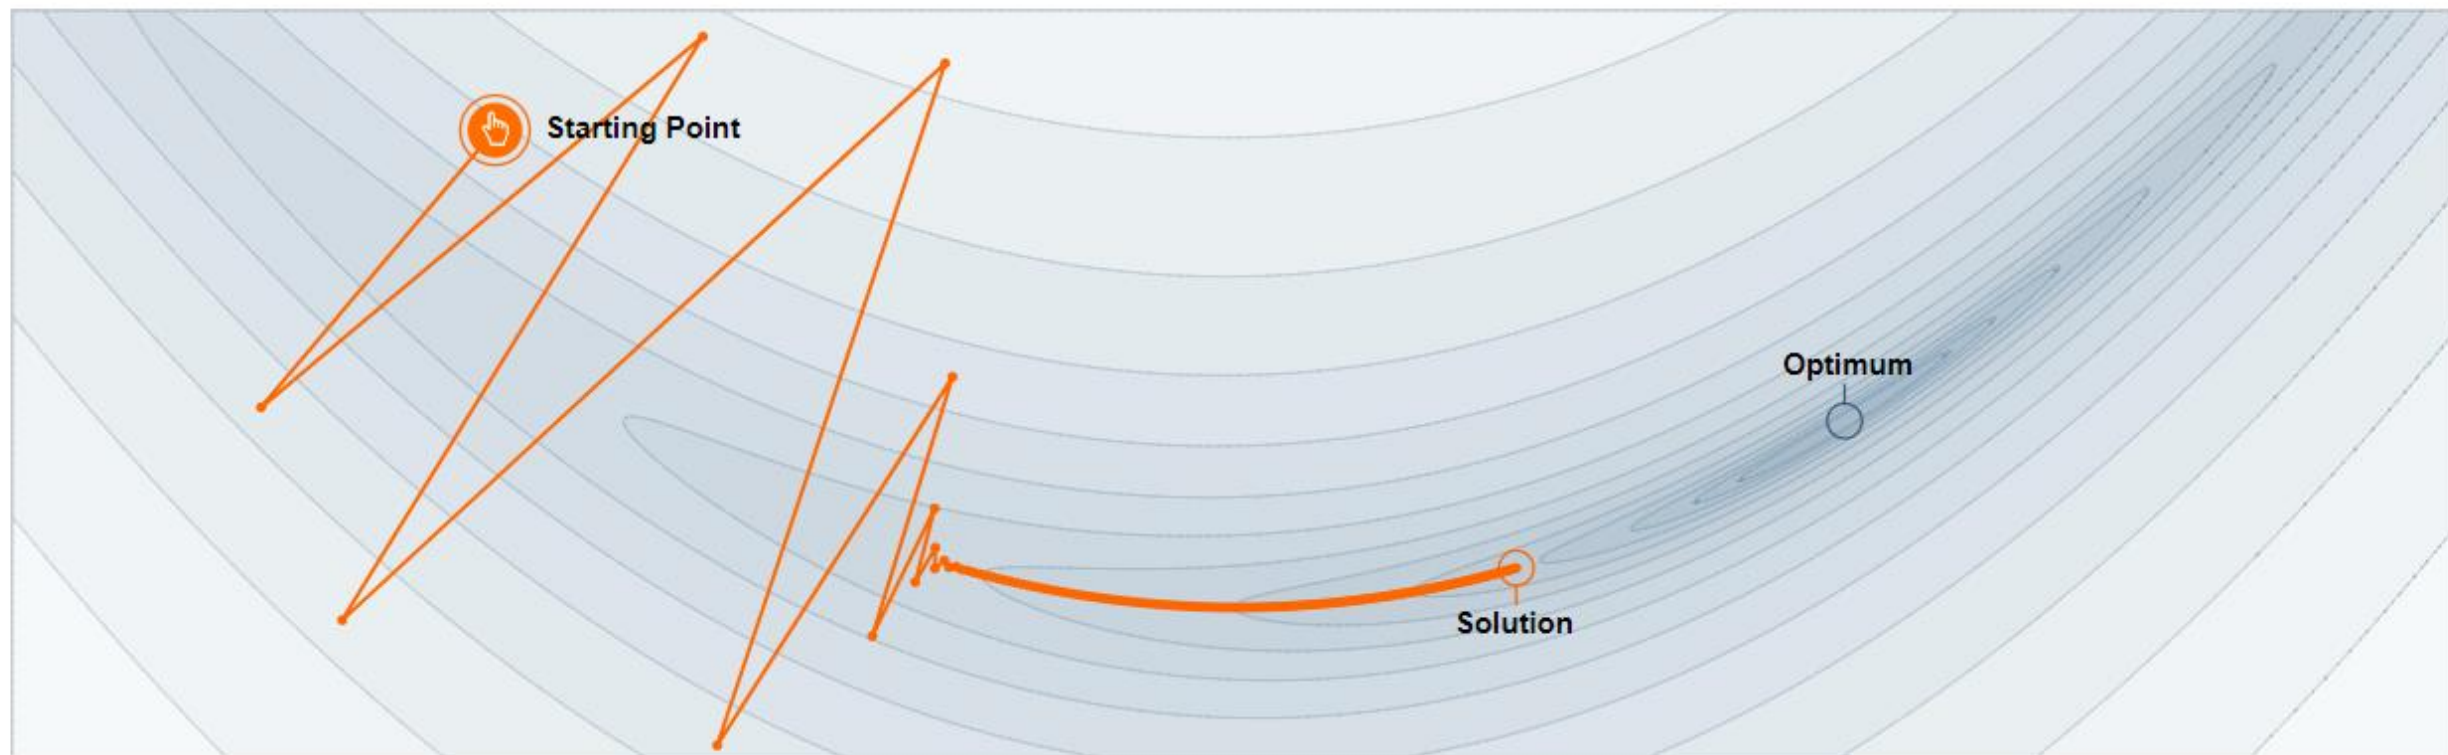
Momentum

$$w^{k+1} = w^k - \alpha \nabla f(w^k).$$

$$z^{k+1} = \beta z^k + \nabla f(w^k)$$

$$w^{k+1} = w^k - \alpha z^{k+1}$$

No momentum



Step-size $\alpha = 0.0038$

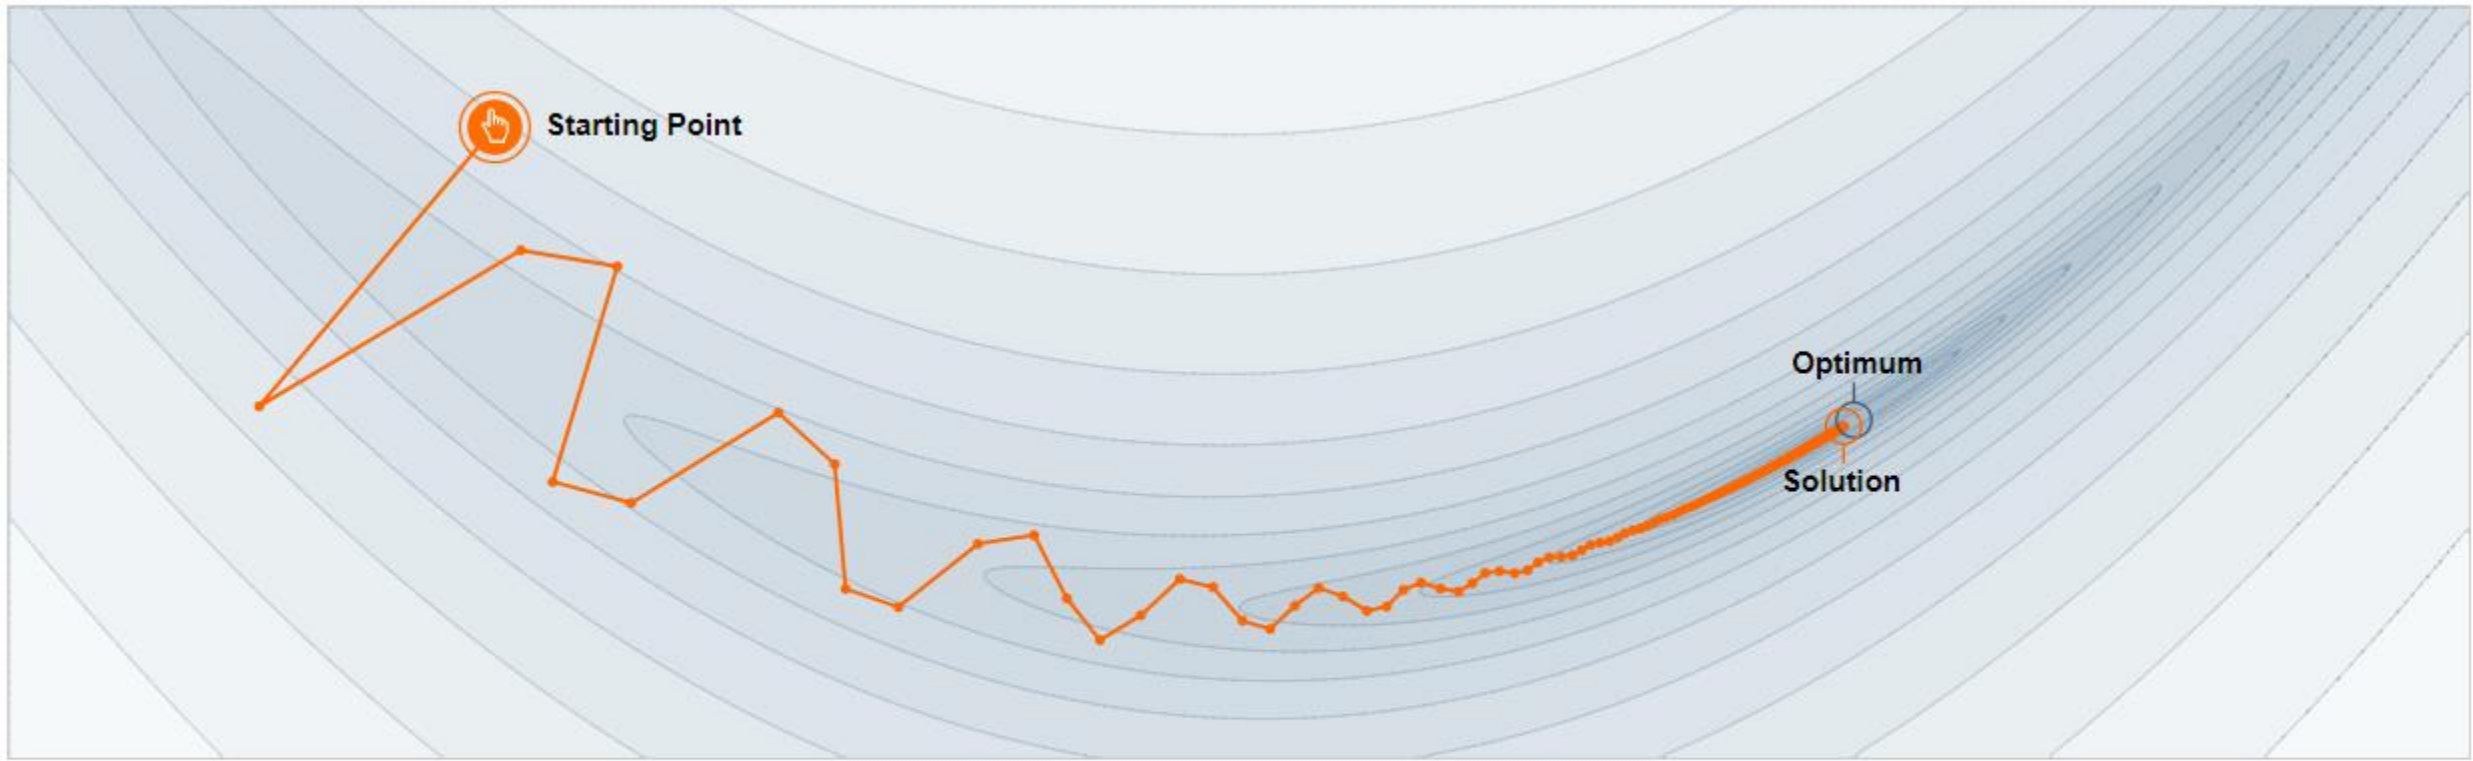


Momentum $\beta = 0.0$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

With appr. momentum



Step-size $\alpha = 0.0038$



Momentum $\beta = 0.78$




We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

momentum



```
optimizer = torch.optim.SGD(model.parameters(), args.lr,  
                             momentum=args.momentum,  
                             weight_decay=args.weight_decay)  
scheduler = ReduceLROnPlateau(optimizer, 'min')  
  
for epoch in xrange(args.start_epoch, args.epochs):  
    train(train_loader, model, criterion, optimizer, epoch)  
    result_avg, loss_val = validate(val_loader, model, criterion, epoch)  
    scheduler.step(loss_val)
```

Learning rate tuning





Andrej Karpathy 

@karpathy



3e-4 is the best learning rate for Adam, hands down.

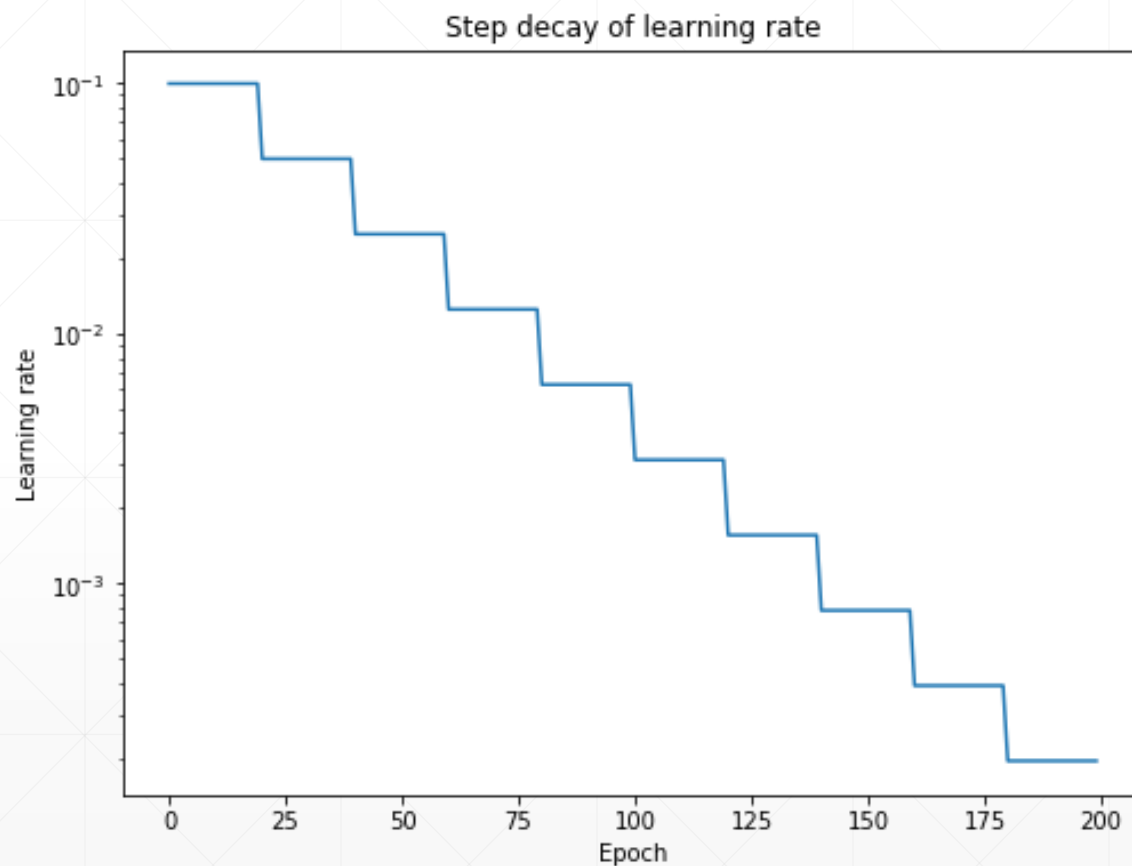
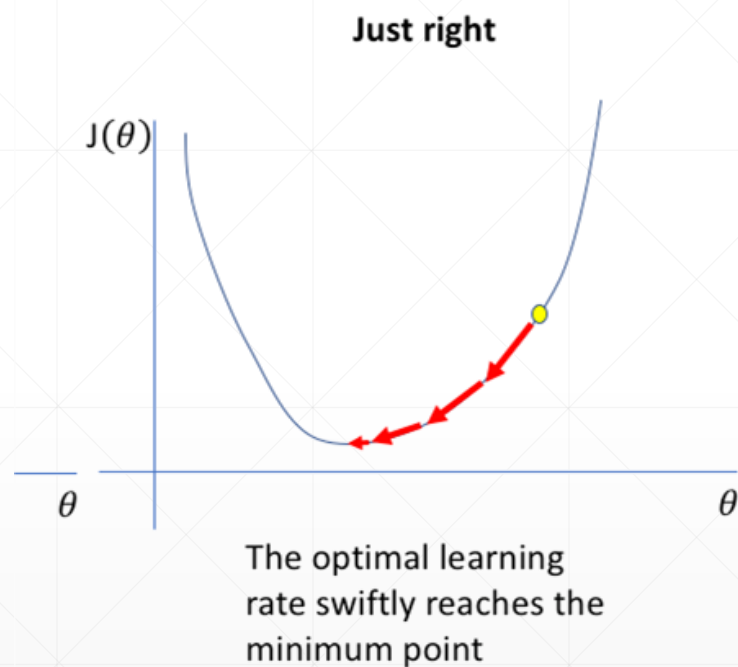
♡ 408 11:01 AM - Nov 24, 2016

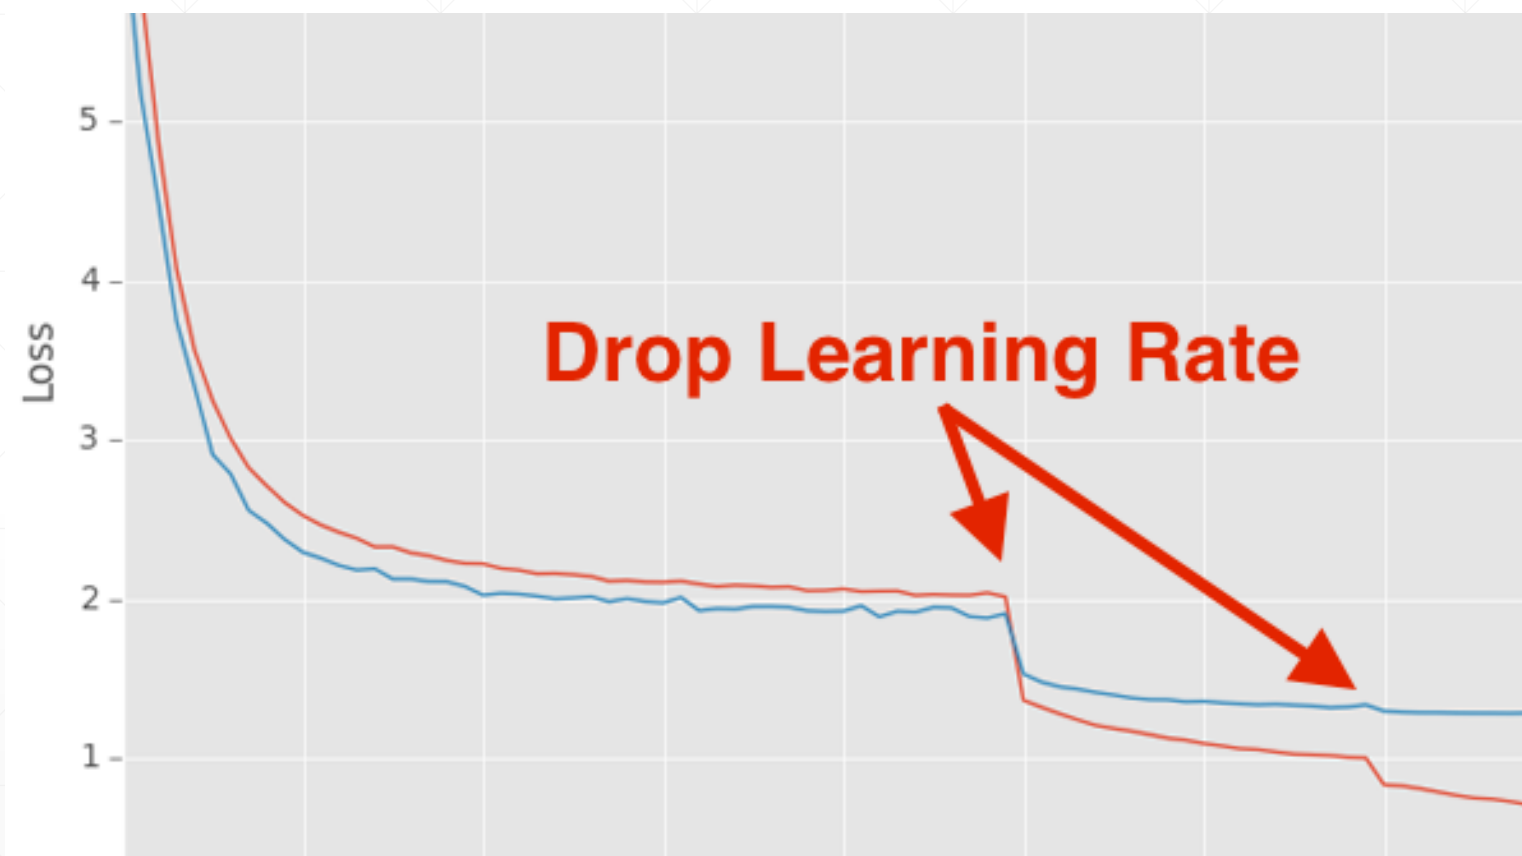


💬 124 people are talking about this



Learning rate decay





Scheme 1.

CLASS `torch.optim.lr_scheduler.ReduceLR0nPlateau(optimizer, mode='min', factor=0.1, patience=10, verbose=False, threshold=0.0001, threshold_mode='rel', cooldown=0, min_lr=0, eps=1e-08)`

[SOURCE]



```
optimizer = torch.optim.SGD(model.parameters(), args.lr,
                             momentum=args.momentum,
                             weight_decay=args.weight_decay)
scheduler = ReduceLR0nPlateau(optimizer, 'min')

for epoch in xrange(args.start_epoch, args.epochs):
    train(train_loader, model, criterion, optimizer, epoch)
    result_avg, loss_val = validate(val_loader, model, criterion, epoch)
    scheduler.step(loss_val)
```

Scheme 2.



```
# Assuming optimizer uses lr = 0.05 for all groups
# lr = 0.05      if epoch < 30
# lr = 0.005     if 30 <= epoch < 60
# lr = 0.0005    if 60 <= epoch < 90
# ...
scheduler = StepLR(optimizer, step_size=30, gamma=0.1)
for epoch in range(100):
    scheduler.step()
    train(...)
    validate(...)
```

下一课时

其他Tricks

Thank You.
